

# Typhoon: A Slice-Scrambled In-Place LSD Sort

Zelun Liu  
Texas A&M University  
zelunliu@tamu.edu

Arif Arman  
Texas A&M University  
arman@tamu.edu

Dmitri Loguinov  
Texas A&M University  
dmitri@cse.tamu.edu

**Abstract**—This paper designs a novel in-place LSD (least significant digit first) radix sort for data-intensive applications. Our framework, which we call *Typhoon*, drops the histogram pass on each level of the sort except the last one, incorporates a high-performance architecture for dynamically expanding output buckets using low-overhead memory blocks we call *slices*, and includes a number of optimizations that reduce pipeline stalls due to cache conflicts and read-after-write bottlenecks. Because Typhoon scatters slices in each bucket randomly across RAM, it has to employ novel mechanisms for non-linear prefetch that combat tendencies of the CPU to pollute the cache on each jump. At the end of the sort, Typhoon uses OS virtual-memory primitives to unscramble the slices and put them in correct order within the input buffer. Results show that Typhoon achieves a significant single/multi-core improvement over the existing methods, including recent AVX-512 efforts from Google [13] and Intel [17], often doubling or tripling their performance.

## I. INTRODUCTION

Sorting has become a ubiquitous building block behind many big-data computational frameworks and distributed systems, including various MapReduce platforms [1], [2], [9], large-scale databases [5], and external-memory graph analytics [8]. After decades of research, improving sort performance has become a difficult target, which we formalize in *five main parameters* – single-threaded speed, robustness against adversarial/non-uniform inputs, RAM usage during the sort (i.e., in-place vs out-of-place), stability (i.e., preservation of original order between duplicate keys), and performance in multi-core environments.

In particular, single-threaded speed measures an algorithm’s useful work per CPU cycle, which is a top priority in scenarios that are not bottlenecked by the total RAM bandwidth of the system. This may include HBM (High Bandwidth Memory) server architectures (e.g., 900 GB/s per socket [12]), optimizing for power consumption, and/or sorting on fewer than all available cores. Resilience against non-uniformity guarantees predictably high performance on real-world datasets, which are often skewed, while in-place operation either saves on hardware cost (i.e., requires half the RAM) or allows fewer passes in external memory compared to out-of-place methods. Stability is an important property in key-value sorts (e.g., in databases and MapReduce), where it is crucial to ensure that an existing order of values is not disturbed by subsequent sorts of the data. Finally, scaling behavior in multi-core settings reflects the algorithm’s synchronization overhead and combined memory traffic across multiple threads, which in some cases can become a separate choke point.

Unfortunately, prior work exhibits a tradeoff between these objectives, which includes sensitivity to key distribution [6],

**Algorithm 1:** Textbook LSD

```
1 Func LSD(Item *input, int n)
2     allocate aux array of size n
3     for (L = 0; L < ⌈w/b⌉; L++) do
4         if (L & 1) == 0 then
5             Split(input, n, aux, L);    ▷ even level
6         else
7             Split(aux, n, input, L);  ▷ odd level
8
9 Func Split(Item *in, int n, Item *out, int L)
10    buck = Histogram(in, out, L);   ▷ set up pointers in out array
11    for (i = 0; i < n; i++) do
12        idx = ExtractIdx(in[i], L);  ▷ bucket index
13        *buck[idx]++ = in[i];      ▷ write item, increment pointer
```

[15], [19], out-of-place and/or unstable operation [3], [6], [13], [15], [19], [30], [31], low speed [4], [7], [25], and non-trivial complexity in achieving efficient multi-threading [6], [13], [15]. Our goal in this paper is to develop a sorting framework that not only rivals the existing methods in terms of robustness, stability, and RAM usage, but also surpasses them in single/multi-core performance.

Due to limited space, the presentation in this paper is brief; however, a more detailed discussion, additional benchmarks, and deeper analysis can be found in the longer version of the paper [20].

## II. STATIC TYPHOON (S-TYPHOON)

### A. Baseline LSD

Assume  $n$  input keys, each consisting of  $w$  bits. Algorithm 1 shows a textbook out-of-place version of LSD (least significant digit first) radix sort, which is stable by design. After creating an auxiliary array of size  $n$  (Line 2), the method runs  $\lceil w/b \rceil$  passes (levels) that alternate between partitioning the input into the aux buffer and vice versa (Lines 3-7), where  $b \geq 1$  is the number of bits examined at each level. The `Split` function begins with a histogram (Line 10) that sets up  $2^b$  destination bucket pointers `buck`, where `buck[0] = out` and `buck[i+1] - buck[i]` is the number of keys that will be written into the  $i$ -th bucket. Following this, Line 12 isolates the  $b$  bits that represent the bucket index of each key `in[i]` and Line 13 writes the item into the corresponding memory location, updating the destination pointer `buck[idx]` in the process.

On each level, Algorithm 1 reads  $n$  items from RAM during the histogram pass and another  $n$  during splitting. On top of that, it writes  $n$  items to the output buckets, which causes the CPU to additionally read for ownership all destination cache lines in the output buffer. Thus, Algorithm 1 ends up with

---

**Algorithm 2: WCv1**


---

```

1 Func Split(Item *in, int n, Item *out, int L)
2   buck = Histogram(in, out, L);
3   for (i=0; i < n; i++) do
4     prefetch (in + i + D);
5     idx = ExtractIdx(in[i], L);
6     p = tmpBuckets + idx*B;
7     p[tmpSize[idx]] = in[i];
8     if ++tmpSize[idx] == B then
9       OffloadAVX(buck[idx], p);
10      buck[idx] += B;
11      tmpSize[idx] = 0;
12
13 Func OffloadAVX(_m256i *dest, _m256i *src)
14   for (i=0; i < R / sizeof(_m256i); i++) do
15     x = _mm256_load_si256(src + i);
16     _mm256_stream_si256(dest + i, x);

```

---

TABLE I  
WCv1A SPEED

run	len	M/sec	c/key
	1	1,121	4.2
	4	938	5.0
	16	826	5.7
	512	883	5.3

a total of  $4n$  keys of memory traffic per level. One can do significantly better by utilizing software *write-combine* (WC) [5], [15], [19], [25], [27], [28], [33], which initially stores data into small tmp buckets contained in the L1/L2 cache and then offloads them to RAM using non-temporal (streaming) stores that bypass the cache. For efficiency reasons, tmp bucket size  $B$  is usually assumed to be a multiple of cache-line size. Because streaming avoids read-for-ownership, this reduces RAM traffic to  $3n$  per level, i.e., 25% lower than the naive approach. Assuming  $2^b$  is larger than the TLB size, which is commonly the case, this also decreases the number of TLB misses from one per key in Algorithm 1 to one per  $B$  keys.

Drawing inspiration from [15], which is currently the fastest implementation of WC, assume tmpBuckets is an array of  $(2^b \cdot B)$  items and tmpSize[i] stores the current number of keys in bucket  $i$ . Algorithm 2, which we call WCv1, shows the baseline partitioning function of an optimized LSD. Its Line 4 runs a prefetch at some distance  $D$  in the input buffer, Line 8 detects tmp bucket overflow, and Line 9 uses a non-temporal memcpy from the start of the tmp bucket (i.e., pointer  $p$ ) to the corresponding location in RAM (i.e.,  $buck[idx]$ ). Note that  $R = B \cdot \text{ItemSize}$  in Line 14 is the length of each bucket in bytes, which is assumed to be a multiple of 256-bit AVX register data type *\_m256i*.

To make discussion more focused, it should be noted that  $b = 8$  is currently the optimal value for LSD, in both Algorithm 1 and 2. Larger values of  $b$  (such as 10 or 11) may be used to reduce the number of passes; however, each pass becomes significantly slower due to the increased number of TLB misses and page-table walks, making performance of the whole sort noticeably worse. Therefore, most of the examples below assume  $b = 8$  and 256-way partitioning. In this setup, sorting 32-bit keys (i.e., four levels) requires  $16n$  memory traffic in Algorithm 1 and  $12n$  in Algorithm 2.

### B. Overview of S-Typhoon

While the basic partitioning engine in Algorithm 2 is a good starting point, we are interested in the question of achieving the absolute maximum performance, both in terms of CPU cycles per key and RAM traffic, in order to establish a definitive

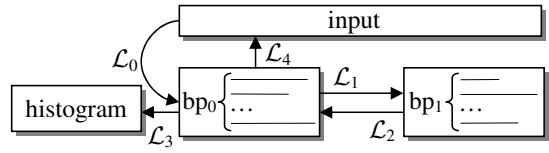


Fig. 1. S-Typhoon workflow overview (32-bit keys).

upper bound on LSD speed. At a minimum, each pass of the sort has to read through the input keys, decide on their buckets, and send them to output, i.e., the work done in Lines 3-11 must remain. However, the histogram in Line 2 can be omitted unless keys in adjacent output buckets must appear contiguously in RAM, which is a requirement only for the final level. Avoiding the histogram not only reduces the amount of CPU cycles by  $\sim 30\%$ , but also leads to lower RAM pressure in multi-threaded scenarios. If this idea can be implemented efficiently in practice, 32-bit keys would require three levels of  $2n$  memory traffic (i.e., non-counting) and one level of  $3n$  (i.e., counting), for a total of  $9n$  per full sort, dropping the memory load of Algorithm 2 by another 25%.

Eliminating the histogram requires expanding buckets dynamically without knowing their final size. We delay this issue until the next section, but in the meantime assume that an oracle pre-allocates two sets of static buckets,  $bp[0]$  and  $bp[1]$ , that never overflow. In this notation,  $bp[L&1][i]$  points to bucket  $i$  during level  $L$ . Then, Static Typhoon (S-Typhoon) proceeds in five steps shown in Fig. 1, where partitioning levels  $\mathcal{L}_0 - \mathcal{L}_2$  are non-counting, level  $\mathcal{L}_3$  is the histogram on the most-significant digit, and the final distribution pass  $\mathcal{L}_4$  returns all items back into the input buffer. Note that buckets during  $\mathcal{L}_0 - \mathcal{L}_2$  are kept in disjoint memory locations, while those in  $\mathcal{L}_4$  are tightly packed (i.e., appear with no gaps).

### C. Read-After-Write Bottlenecks

Our first topic is to analyze WCv1 under skewed key distributions. While the method works well for uniform keys, our results show that it inexplicably reduces speed when multiple adjacent keys are sent into the same bucket, which is a common occurrence for non-uniform inputs. Define WCv1a to be Algorithm 2 without the histogram pass in Line 2. Then, Table I demonstrates this issue with 1 GB of 32-bit keys and an Intel Skylake-X i7-7820X clocked at a fixed 4.7 GHz, where the run length in the first column specifies how many back-to-back copies of each random key are generated. The speed begins at 1,121M keys/sec (4.2 cycles/key) in the first row (all unique), drops to 826M/sec in third row (5.7c/key), and finally settles on 883M/sec (5.3c/key) for sufficiently long runs. The worst case is 1.4 $\times$  slower than the best, which is not a negligible drop.

A key component of data partitioning, whether it uses write-combine or not, is to perform updates to shared counters as items arrive from input. These can be bucket pointers in Algorithm 1 (Line 13) or tmp bucket sizes in Algorithm 2 (Line 8). To better understand the challenges the CPU faces in these cases, first consider a simplified problem, also

**Algorithm 3:** Histogram Hv1

```

1 Func Hist(Item *in, int n)
2   for (i=0; i < n; i++) do
3     prefetch (in + i + D);
4     idx = *(uint8*)(in + i);
5     hist[idx]++;

```

TABLE II  
Hv1 SPEED

run len	M/sec	c/key
1	2,250	2.1
4	1,817	2.6
16	1,454	3.2
512	927	5.1

of relevance to S-Typhoon, whose purpose is to compute a histogram using the first byte of each item in an array. A baseline solution [19], [25], which we call Hv1, is illustrated by Algorithm 3. Its performance using the same setup as before is shown in Table II. In this case, the reduction in speed is even sharper, i.e., 2.4× between all-unique and all-duplicate.

To delve deeper, consider the CPU pipeline for Hv1, which repeats a pattern of two loads, an increment, and a store:

```

idx0 = *(uint8*)in; c0 = hist[idx0]; inc c0; hist[idx0] = c0;
idx1 = *(uint8*)(in+1); c1 = hist[idx1]; inc c1; hist[idx1] = c1;

```

where *idx0*-*idx1*, *c0*-*c1* are registers. For performance reasons, the CPU’s out-of-order execution engine attempts to hoist loads ahead of preceding stores unless its memory-disambiguation module detects a conflict. As the pipeline decodes the uops, it sees a store to *hist[*idx0*]*, followed by a load from *hist[*idx1*]*, both from yet-unknown buckets *idx0*, *idx1*. Without additional hints, the CPU cannot decide whether this presents a conflict and optimistically assumes that these uops are independent, which causes it to reorder the load from *hist[*idx1*]* to precede the store into *hist[*idx0*]*.

When adjacent keys refer to different buckets, load hoisting allows higher levels of instruction-level parallelism as counters from multiple locations can be fetched and incremented concurrently. However, doing the same for pairs of keys with *idx0* = *idx1* leads to consistency violations, which are detected by the CPU just before these instruction retire, causing expensive pipeline flushes. It is speculated [11] that Intel maintains a history of mispredictions for each load and temporarily disables hoisting after a threshold of violations is reached.

While hoisting is inoperable and duplicate keys are still arriving, the histogram runs into dependency chains between each load from the L1 cache and the preceding store. Going into a latency-bound regime (i.e., 4c/load for Skylake-X) is already a major bottleneck, but an additional problem arises from having to search through the store buffer and forward loads out of it [16], which sometimes has an even higher latency [34]. Breaking loop-carried dependencies and reducing frequency of pipeline stalls is our next topic.

#### D. Reducing Store-Forwarding Costs

This analysis gives rise to the following idea. Dependency between *k* adjacent store-load pairs can be resolved by reading *k* histogram counters upfront, performing comparison across all *k(k* − 1)/2 pairs of bucket indexes, and incrementing the relevant counters using the result of the comparison. For *k* = 2, this is demonstrated by Algorithm 4, which simultaneously reads two indexes from input and obtains their

**Algorithm 4:** Histogram Hv2

```

1 Func Hist(Item *in, int n)
2   for (x = in; x < in + n; x += 2) do
3     prefetch (x + D);
4     idx0 = *(uint8*)x;
5     idx1 = *(uint8*)(x + 1);
6     c0 = hist[idx0]; c1 = hist[idx1];
7     c0++; hist[idx0] = c0;
8     c1 = (idx0 == idx1) ? c0 : c1;
9     c1++; hist[idx1] = c1;

```

TABLE III  
Hv2 SPEED

run len	M/sec	c/key
1	2,496	1.9
4	2,478	1.9
16	2,275	2.1
512	1,402	3.4

counters at the start of each iteration (Lines 4-6), keeping them in registers. It then updates counter *c0* (Line 7) and decides the value of *c1* based on whether the two buckets are the same (Line 8). To keep the algorithm branchless, ternary operator ? is implemented using a *conditional move* CPU instruction *cmov*, which needs only 0.5c/iteration (i.e., 0.25c/key). If the compiler does not issue *cmov*, assembly can be used instead.

Table III displays the performance of Hv2. For unique keys in the first row, the speed goes up 11% compared to Hv1. As the run length increases, performance gains become more substantial, where Hv2 finishes with a 1.5× advantage in the last row compared to Table II. The gradual reduction in speed as the burst length increases in Table III can be explained by Hv2’s success at eliminating store-load dependencies *within* each pair of items, but not *between* pairs. Unfortunately, for *k* ≥ 3, the quadratic cost of doing all-to-all scalar comparisons becomes prohibitively high, i.e., no further improvement is currently possible under the same umbrella.

We therefore leverage conflict detection in Algorithm 4 by applying the same principles to key distribution. This is shown in Algorithm 5 under the name of WCv2. Unlike the previous version WCv1a, bucket addresses in RAM are passed into the function in the third argument (i.e., array of pointers *buck*), which refers to either *bp[0]* or *bp[1]* depending on the level. Additionally, the fourth argument specifies an array of pointers *t* to the starting position in each tmp bucket. The algorithm further assumes that the number of bytes in each tmp bucket *R* = *B* · *ItemSize* is a power of two and the start of *tmpBuckets* is aligned to *R* bytes.

For each pair of items, macro *MOVE* reads their indexes from the corresponding byte of the key (Lines 7-8) and obtains both pointers *p0*, *p1* (Line 9) without yet knowing if there exists a conflict. It then calls a macro *WRITE*, which stores the first key into *p0* and checks for the end of the tmp bucket (Line 16). If so, it jumps *p0* to the start of that bucket and offloads its *B* items to RAM. At the end, it updates *t[*idx0*]* in Line 20 to reflect the new position. At this point, *p1* can be computed using a conditional move in Line 11 and the process repeats for the second key. Note that WCv2 avoids store forwarding within each pair of keys because *t[*idx1*]* is read (Line 9) before *t[*idx0*]* is written (Line 20).

Table IV shows the resulting speed. Compared to WCv1a in Table I, this version gains a few M/sec in the first row, improves by 20% in the second, 35% in the third, and 47% in the last one.

**Algorithm 5:** WCv2

```

1 Func Split(Item *in, int n, Item **buck, Item **t, int L)
2   for (x = in; x < in + n; x += 2) do
3     | prefetch (x + D);
4     | MOVE(x);
5
6 Macro MOVE(x)
7   idx0 = *((uint8*)x+L);
8   idx1 = *((uint8*)(x+1)+L);
9   p0 = t[idx0]; p1 = t[idx1];
10  WRITE(x[0], p0, idx0);
11  p1 = (idx0 == idx1) ? p0 : p1;
12  WRITE(x[1], p1, idx1);
13
14 Macro WRITE(key, p, idx)
15   *p++ = key; > store item
16   if (p & (R-1) == 0) then > overflow?
17     | p -= B; > roll back to start of bucket
18     | OffloadAVX(buck[idx], p);
19     | buck[idx] += B;
20     | t[idx] = p;

```

TABLE IV  
WCv2 SPEED

run	len	M/sec	c/key
1	1,128	4.2	
4	1,128	4.2	
16	1,118	4.2	
512	1,302	3.6	

### E. Histograms Revisited

We now deal with the issue of designing the histogram for level  $\mathcal{L}_3$ , keeping in mind that we no longer need this solution to be applicable to a key splitter. Our next approach, illustrated as Hv3 in Algorithm 6, unrolls the loop to grab multiple keys at once and writes updated counters into *separate* histograms. While this example shows unrolling to  $r = 4$  keys and  $h = 4$  histograms, other combinations are possible as well. Furthermore, when  $r$  exhausts general-purpose registers, our implementation of Hv3 uses SIMD (SSE/AVX) to hold the keys. A crucial element of this technique is the use of an *offset*, which specifies the distance (in bytes) between the start of each histogram and the end of the previous one. Since each histogram is exactly  $256 \times 8 = 2$  KB, offsets are needed to avoid 4K aliasing and conflicts in set-associative caches [16].

Using  $r = 16$  keys and  $h = 8$  histograms, the upper half of Table V shows that a zero offset can produce a 16-35% improvement over Hv2, but the resulting method still chokes on duplicate keys, losing over a billion keys/sec between the first and last rows. On the other hand, offsetting the histograms by 8 bytes yields a drastically different result, i.e., a constant 1.6c/key, as also shown in the table.

### F. Multi-Threading

Assume a joint sort across  $T$  threads, each holding its own dual set of RAM buckets  $\text{bp}[0]$ ,  $\text{bp}[1]$  and local tmp buckets in the corresponding L1/L2 cache. Suppose matrix  $M$  consists of all sub-buckets written by the threads after a particular level of splitting, i.e.,  $M_{ij}$  represents the contents of bucket  $i$  created by thread  $j$ . To identify a sub-bucket, it is sufficient to specify its 2D index  $(ij)$ . Furthermore, let triple  $(ijr)$  refer to the  $r$ -th key in bucket  $M_{ij}$ . Then, a *row-major order on keys* is defined as  $(xyr) \prec (uvt)$  iff  $(x < u) \vee (x = u, y < v) \vee (x = u, y = v, r < t)$ . Note that for level  $\mathcal{L}_0$ , we assume  $M = (M_{00})$  is a  $1 \times 1$  matrix consisting of the input array.

Multi-threading requires assigning each thread  $p = 1, 2, \dots, T$  a set of keys  $\Delta_p$  consisting of  $n/T$  triples  $(ijr)$

**Algorithm 6:** Histogram Hv3

```

1 Func Hist(Item *in, int n)
2   for (x = in; x < in+n; x += 4) do
3     | prefetch (x + D);
4     | idx0 = *((uint8*)x);
5     | idx1 = *((uint8*)(x+1));
6     | idx2 = *((uint8*)(x+2));
7     | idx3 = *((uint8*)(x+3));
8     | hist0[idx0]++;
9     | hist1[idx1]++;
10    | hist2[idx2]++;
11    | hist3[idx3]++;

```

TABLE V  
Hv3 SPEED

run	len	M/sec	c/key
offset = 0			
1	2,912	1.6	
4	2,688	1.7	
16	2,215	2.1	
512	1,904	2.5	
offset = 8			
1	2,941	1.6	
4	2,941	1.6	
16	2,941	1.6	
512	2,941	1.6	

such that  $\{\Delta_1, \dots, \Delta_T\}$  forms a partition on  $M$ . Two rules must be satisfied in order to ensure correctness: a) each thread processes keys assigned to it in row-major order; and b) if  $(xyr) \in \Delta_p$  and  $(uvt) \in \Delta_q$ , where  $p < q$ , then  $(xyr) \prec (uvt)$  must hold. This guarantees stability, i.e., that the next level of LSD does not break the relative order established within each bucket on the previous level. To achieve b), our load-balancing algorithm views all  $n$  keys as a one-dimensional array in the row-major order of  $M$  and assigns its  $p$ -th consecutive batch of size  $n/T$  to thread  $p$ .

The cost of computing these boundaries depends on the number of buckets  $2^b$  and thread count  $T$ , but this is usually negligible (i.e., under 10  $\mu$ s) compared to the sort time. Since threads require data in sub-buckets  $M_{ij}$  created by other threads, each level  $\mathcal{L}_0 - \mathcal{L}_4$  ends with a barrier that synchronizes the threads.

### III. TYPHOON

To make the framework developed in the previous section practical, the first challenge is to design a low-overhead technique for dynamically expanding output buckets as they are written to. While this problem has been touched upon in prior work [4], [15], [25], [30], performance of these solutions leaves much to be desired. Our objective here is to create a new bucket-management infrastructure that runs  $\mathcal{L}_0 - \mathcal{L}_4$  at almost the same speed as S-Typhoon, but without using an oracle to statically pre-allocate the buckets. After  $\mathcal{L}_4$ , the data will end up in a number of disjoint locations in RAM, where the second challenge is to restore proper order between the keys using a novel unscrambling level we call  $\mathcal{L}_5$ .

Define a *slice* to be contiguous region of  $S$  bytes in virtual memory starting from an address that is aligned to  $S$ . Slices come from two places – the input buffer and some auxiliary space that is needed to provide support to partially filled slices during the split. To speed up detection of end-of-slice, we assume  $S$  is a power of 2, and to prevent offloads from crossing slice boundaries, let  $S$  be a multiple of tmp bucket size  $R$ . For reshuffling at  $\mathcal{L}_5$ ,  $S$  must also be a multiple of page size.

#### A. Data Structures

We start by considering single-threaded execution. Suppose the sort maintains a stack of free slices, which is an array of 64-bit pointers to the start of each slice. Compared to other data structures, stacks have an advantage in their low push/pop

**Algorithm 7:** WCv4 (simplified)

---

```

1 Func Split(Item **s, int ns, Item **bp, Item **t, int L)
2   for (j=0; j < ns; j++) do    ▷ iterate over all slices
3     cur = s[j]; next = s[j+1]; dist = next - cur;
4     do
5       prefetch (cur + dist);    ▷ prefetch next slice
6       MOVE(cur);
7       cur += 2;
8     while cur & (sliceSize - 1);

```

---

cost, i.e., one `stackTail` pointer, a load/store instruction, and a register increment/decrement. Additionally, stacks achieve high temporal locality because of immediate reuse of slices between input and output. As we see below, this allows Typhoon to run certain levels of the sort faster than S-Typhoon. At the start of the sort, the free stack is initialized to  $A$  auxiliary slices, where  $A$  determines the  $O(1)$  constant in  $n + O(1)$  memory usage. Each level of the sort requires at least  $2^b$  slices in the free stack. Thus,  $A \geq 2^b$  must hold. We refine this bound later in the section.

To keep track of the slices assigned to each bucket, suppose *slice database* *sd* contains in *sd.p[i][j]* the address of the  $j$ -th slice in bucket  $i$ . For each new slice popped from the free stack, the splitter records tuple (bucket  $\text{idx}$ , slice pointer) into a separate pre-allocated buffer. After the level is over, this array is processed to count the number of slices that went into each bucket  $i$ , which allows easy construction of *sd*. Since Typhoon alternates between two sets of buckets, each set requires a separate slice database, which we call *sd[0]* and *sd[1]*. If the sort reads slices from *sd[k]*, where  $k$  is either 0 or 1, it keeps track of the new ones in *sd[1-k]*.

### B. Aligned Splitter ( $\mathcal{L}_0 - \mathcal{L}_2$ )

We extend the S-Typhoon WCv2 splitter, which we now call WCv3, to accept an array of slice pointers *s[]* rather than one large buffer. The main loop in Algorithm 5 remains essentially the same, except it gets interrupted every  $S$  bytes to load the next input slice pointer. This involves four CPU instructions per slice and a mispredicted branch.

On output, the `OffloadAVX` function has to check if the destination pointer *buck[idx]* is aligned to slice boundary using bitwise masking. If so, a free slice is popped from the stack and added to the array *sd[k].p[idx]*. To ensure proper operation at the start, all buckets begin such that *buck[idx] = NULL*, which causes a trip to the stack on first access to each destination pointer. In total, output slice management requires one mask instruction and a well-predicted branch *per offload*, which adds at most 1/4 cycle per  $R$  bytes, as well as 6 additional mov/add/sub instructions *per slice*, which add  $\sim 1.5$  cycles per  $S$  bytes. Because of the simplicity of its data structures, WCv3 can fit all pointers and variables into 13 general-purpose registers, leaving three unused. This ensures no register pressure, spills to the stack, or reloads.

Considering the low cost of managing the slice database and the stack, it is perhaps unexpected that WCv3 runs a lot slower than WCv2 during  $\mathcal{L}_1$ . With 4-KB slices and 32-bit keys, Intel i7-7820X shows a reduction in speed from

TABLE VI  
LEVEL  $\mathcal{L}_1$  SPLITTER SPEED (M/SEC)

WCv2 (static)	WCv3		WCv4	
	4 KB	8 KB	4 KB	8 KB
1,128	872	939	1,117	1,139

1,128M/sec to 872M/sec, a loss of 23%! Further investigation reveals that this issue is caused by two compounding effects – software prefetch at distance  $D$  in Line 3 of Algorithm 5, which pollutes the cache with irrelevant data ahead of each jump, and CPU hardware prefetchers that detect scans and also load some amount of garbage following each slice.

There is not much we can do about the latter issue, but the former can be alleviated by introducing a non-linear prefetch into WCv3, which keeps both the current slice pointer *cur* = *s[j]* and the next one *s[j+1]* in registers, prefetching at address  $(x + D)$  when  $x + D < cur + S$  and  $(next + x - cur + D - S)$  otherwise. In cases when the CPU allows a range of prefetch options sufficiently-far in the future to work at optimal speed, which is the case for Intel and AMD, the loop can be further simplified to always prefetch from the next slice. For these situations, Algorithm 7 shows a high-level operation of the new approach WCv4.

Generally, it is expected that larger slices are faster because the CPU prefetches less garbage compared to the amount of useful data in the slice. Additionally, the cost of managing the free stack and database *sd* becomes smaller as well. Results in Table VI confirm this observation using 4 and 8 KB slices. Even with 8-KB slices, WCv3 struggles to match the static speed. On the other hand, WCv4 essentially ties S-Typhoon with 4KB slices and exceeds its performance using 8KB slices. As mentioned earlier, stack-based reuse of slices sometimes gives Typhoon an advantage over S-Typhoon.

### C. Histogram ( $\mathcal{L}_3$ )

As level  $\mathcal{L}_3$  runs almost 3× faster than the splitter, the effect of incorrect prefetch becomes even worse. In particular, the speed of Hv3 drops by 47%, i.e., from 2,941M/sec in Table V to 1,560M/sec (4-KB slices). Applying the trick from Algorithm 7 improves the result to 2,744M/sec, but this is still 200M/sec slower than static.

In contrast to the splitter, which must process the items in row-major order within matrix  $M$ , the histogram is not constrained in how it visits the keys. Hence, it is sufficient to identify all *contiguous* runs of keys in virtual memory and call Hv3 on each of them. To determine these regions, one option is to sort all slice pointers in the *sd* database and merge the ones next to each other. However, sorting 256K slices per GB of data is too expensive, which makes the result slower than random jumping using the prefetch of WCv4.

On the other hand, the same outcome can be obtained by extracting the chunks of *unused* space from the free stack and *partial* slices in *sd*, sorting them, and computing the contiguous runs of valid keys by complementing the empty space. Note that partial slices combined with the empty stack

amounts to at most  $A$  separate regions of free memory. Thus, the sort involves a few hundred integers, regardless of  $n$ . We call this method **Hv4** and note that it easily hits the speed of S-Typhoon even with 4-KB slices.

#### D. Unaligned Splitter ( $\mathcal{L}_4$ )

For  $\mathcal{L}_4$ , a new challenge arises due to the possibility that function `OffloadAVX` may cross slice boundaries. The general structure for this level follows **WCv4**, except `OffloadAVX` needs to detect when the destination pointer moves to the next slice. In particular, it examines if the current slice can accommodate another  $R$  bytes; if so, it runs the standard (uninterrupted) offload loop. Otherwise, it moves enough keys to finish the current slice, obtains a new one from the stack, and completes the offload there, in both cases using AVX. The frequency of taking the slice-crossing branch is determined by the ratio of tmp bucket size to slice length, i.e.,  $R/S$ . With  $S = 4$  KB, the probability to take the slower branch on Intel CPUs is 6.25% (i.e.,  $R = 256$ ). Combined with the extra cost of branch misprediction, this explains why Typhoon’s  $\mathcal{L}_4$  runs slightly slower than its  $\mathcal{L}_1 - \mathcal{L}_2$ .

#### E. Multi-Threading

We now deal with slice management across  $T$  threads. Assume each of them maintains a local state consisting of a free stack, a tmp bucket buffer, two sets of bucket pointers `bp[0], bp[1]`, and two slice databases `sd[0], sd[1]`. The goal of this setup is to make threads run with as little interaction with each other as possible. To handle partitioning of matrix  $M$ , we extend the row-major order introduced earlier to organize keys by (bucket, thread, slice). For speed reasons, threads never share slices from each sub-bucket, but the rest of the load-balancing algorithm (Section II-F) remains the same.

For  $\mathcal{L}_0 - \mathcal{L}_2$ , if a thread begins a level with  $A_0 = 2^b$  free slices in its local stack, it has sufficient extra memory to leave one almost-empty slice at the end of each bucket. Thus, **WCv4** can run independently and without modification within each thread using its local stack. However, because a thread reads other thread’s sub-buckets, its stack size at the end of a level can be anywhere from zero to the total number of partial slices in the sub-buckets it visited. Thus, each level introduces a size-imbalance into thread stacks, which, if left uncorrected, eventually leads to a crash.

To address this problem, Typhoon introduces a global stack that contains the remaining slices not currently assigned to any thread. After finishing a level, but before the barrier, each thread returns excess slices (i.e., those above  $A_0$ ) to the global stack. Similarly, when a level begins (i.e., after the barrier), each thread acquires the missing slices to bring its local stack size back to  $A_0$ .

This works well for  $\mathcal{L}_0 - \mathcal{L}_2$ , but additional difficulties arise during  $\mathcal{L}_4$ . Because keys must be contiguous in space at the end of the sort, there can be up to  $T \cdot 2^b$  slices that are shared across bucket boundaries, including between different threads. The main challenge here is to prevent allocation of redundant slices when **WCv4** moves into the last (partial) slice of each

bucket. To address this, the last thread of Typhoon that reaches the barrier at the end of  $\mathcal{L}_3$  pre-allocates all shared slices and then kicks off  $\mathcal{L}_4$ . This entails examining  $T \cdot 2^b$  boundaries between sub-buckets, assigning each a slice from the global free stack, and notifying the threads that these slices have been pre-allocated.

The notification is done through the `sd` database – any time **WCv4** aims to obtain a new slice, it loads the pointer `x` for the current slice  $j$  of bucket  $i$  for its thread  $p$ ; if this value is `NULL`, it gets a new slice from its local stack; otherwise, it uses `x` as a pointer to the pre-allocated slice. Not only that, but `x` specifies the exact location within the slice where sub-bucket ( $ip$ ) begins. Combining this logic with generalized offloads that can move across slices (Section III-D) leads to our final **WCv4L4** splitter in Typhoon.

The worst-case memory usage happens at  $\mathcal{L}_4$ . Observe that threads may hold  $T \cdot A_0$  slices in matrix  $M$  before the level begins, they are given  $T \cdot 2^b$  additional shared slices in pre-allocated memory, and they request  $T \cdot 2^b$  new slices immediately after starting the level. As a result, the smallest number of auxiliary slices  $A$  that the sort needs is  $T \cdot 2^b \cdot 3$ . With 4-KB slices, this leads to 3 MB per thread. Compared to 1 GB/thread worst-case in Vortex [15], this is a major improvement.

It should also be noted that the modified histogram **Hv4** applies only to the single-threaded case since complementing the empty space does not reveal which thread is responsible for which slice. Additionally,  $T \geq 2$  leads to tight interleaving of slices, where the length of contiguous regions assigned to each thread, even if they could be determined efficiently, is often no more than 2 slices, which negates the sought-after benefits. Therefore, we run histogram **Hv4** only for  $T = 1$  thread and **Hv3** otherwise.

#### F. Reshuffle ( $\mathcal{L}_5$ )

After  $\mathcal{L}_4$ , the data is located in  $[n \cdot \text{ItemSize}/S]$  slices, which are randomly scattered in RAM. Assume  $\{v_1, v_2, \dots\}$  is a list of pointers in row-major order within matrix  $M$  recorded in the slice database during  $\mathcal{L}_4$ , ignoring duplicate slices shared across adjacent buckets. Now the remaining task is to make each slice  $v_i$  appear at offset  $iS$  in the input buffer. To accomplish this, the Typhoon constructor obtains a chunk of memory `buf`, big enough to hold both  $n$  input items and  $A$  auxiliary slices, using OS primitives that allow physical pages to be mapped/unmapped within this virtual space. For Windows, this translates into a call to `VirtualAlloc` with the `MEM_PHYSICAL` flag. Typhoon then grabs enough physical pages using `AllocateUserPhysicalPages` and maps them to this buffer via `MapUserPhysicalPages`. The `buf` array is then given to the user to fill in the keys.

Note that the OS provides the PFN (physical frame number) of each allocated page, which Typhoon stores internally in the `pfn` array for later use during remapping. Level  $\mathcal{L}_5$  begins with unmapping all slices in `buf` using  $T$  parallel threads. The information within these pages does not get destroyed, but becomes temporarily inaccessible. After finishing the unmap,

TABLE VII  
TYPHOON SPEED (M/SEC) ON 1 GB INPUT (UNIFORM KEYS)

Level	Single core								All cores								
	32-bit keys				64-bit key-value pairs				32-bit keys				64-bit key-value pairs				
	Static	4KB	8KB	16KB	Static	4KB	8KB	16KB	Static	4KB	8KB	16KB	Static	4KB	8KB	16KB	
0	1,128	1,162	1,182	1,183	815	831	854	865	8,308	8,902	8,944	8,846	4,381	4,504	4,505	4,513	
1	1,110	1,126	1,158	1,161	812	790	820	846	8,289	8,355	8,554	8,575	4,386	4,309	4,413	4,456	
2	1,131	1,134	1,167	1,174	828	794	833	853	8,298	8,379	8,554	8,541	4,391	4,327	4,411	4,460	
3	2,941	2,955	2,933	2,934	2,174	2,037	2,035	2,036	20,831	17,197	17,927	18,286	10,354	9,291	9,896	10,217	
4	1,124	1,121	1,148	1,161	814	788	821	846	8,132	8,129	8,345	8,407	4,352	4,219	4,305	4,344	
0-4	<b>256</b>	<b>259</b>	<b>265</b>	<b>266</b>	<b>187</b>	<b>182</b>	<b>189</b>	<b>193</b>	<b>1,878</b>	<b>1,878</b>	<b>1,919</b>	<b>1,922</b>	<b>990</b>	<b>971</b>	<b>991</b>	<b>1,002</b>	
5	9,323	11,441	14,451		4,682	5,798	7,370		59,005	68,237	75,655		28,718	33,789	38,347		
0-5		252	<b>259</b>	261		175	183	<b>188</b>		1,820	1,866	1,874		939	963	976	

TABLE VIII  
STRONG SCALING OF TYPHOON SPEED (M/SEC) ON 1 GB OF 32-BIT KEYS (16-KB SLICES, UNIFORM KEYS)

Level	1 core	2 cores	3 cores	4 cores	5 cores	6 cores	7 cores	8 cores							
0	1,183	2,340	2.0×	3,533	3.0×	4,680	4.0×	5,824	4.9×	6,978	5.9×	8,041	6.8×	8,846	7.5×
1	1,161	2,312	2.0×	3,462	3.0×	4,605	4.0×	5,714	4.9×	6,813	5.9×	7,788	6.7×	8,575	7.4×
2	1,174	2,340	2.0×	3,490	3.0×	4,628	3.9×	5,752	4.9×	6,835	5.8×	7,822	6.7×	8,541	7.3×
3	2,931	5,256	1.8×	7,747	2.6×	10,163	3.5×	12,479	4.3×	14,651	5.0×	16,643	5.7×	18,286	6.2×
4	1,161	2,327	2.0×	3,456	3.0×	4,609	4.0×	5,733	4.9×	6,827	5.9×	7,750	6.7×	8,407	7.2×
0-4	<b>266</b>	<b>524</b>	<b>2.0×</b>	<b>783</b>	<b>2.9×</b>	<b>1,039</b>	<b>3.9×</b>	<b>1,290</b>	<b>4.9×</b>	<b>1,536</b>	<b>5.8×</b>	<b>1,755</b>	<b>6.6×</b>	<b>1,922</b>	<b>7.2×</b>
5	14,451	28,731	2.0×	39,097	2.7×	49,160	3.4×	58,511	4.0×	65,985	4.6×	67,140	4.6×	75,655	5.2×
0-5	261	515	2.0×	768	2.9×	1,018	3.9×	1,262	4.8×	1,501	5.7×	1,711	6.6×	1,874	7.2×

threads jointly construct the `nextPfn` array that specifies the page frames that need to appear in each position in `buf`.

In more detail, assuming  $P$  is page size in bytes, observe that  $(v[i] - buf)/P$  is the offset in the `pfn` array that contains the  $S/P$  page frames from slice  $v_i$ . Similarly, the offset in the `nextPfn` buffer where  $v_i$  should be mapped to is given by  $i*S/P$ . Therefore, construction of `nextPfn` is a sequence of operations `memcpy(nextPfn + i*S/P, pfn + (v[i]-buf)/P, S/P*sizeof(void*))` for all  $i$ . After the threads are done with unmapping and `memcpy`, they synchronize on a barrier and call `MapUserPhysicalPages` with their assigned portion of `nextPfn`. Another barrier follows, after which Typhoon finishes the sort by swapping `pfn` and `nextPfn` pointers in preparation for the next iteration (if needed). The sort can be called repeatedly any number of times, reusing `buf` and other data structures, without causing allocating of new memory.

#### IV. EVALUATION

Our primary benchmark platform is an Intel i7-7820X, which is an 8-core Skylake-X CPU with a 32-KB L1, 256-KB L2, and 16-MB L3, clocked for these experiments at a fixed 4.7 GHz on each core. We run 32 GB of DDR4-3200 RAM in a quad-channel memory configuration, which yields a peak non-temporal AVX `memcpy` bandwidth of 37 GB/s and a maximum AVX read speed of 86 GB/s across 8 cores. Single-threaded Typhoon is bottlenecked by the splitter's 4.2c/key in Table IV, while the performance of the multi-core version is upper bounded by four `memcpy` passes and one read pass, i.e., 2088M keys/sec for 32-bit items and half of that for 64-bit.

Typhoon, whose source code is available from [32], is compiled in Visual Studio 2019, while prior methods are reported using the best achievable speed among Clang 19, Intel oneAPI C++ 2025 (ICX), and VS 2019. Benchmarks run on

Windows Server 2016 and Ubuntu 24.04.

##### A. Static vs Sliced Typhoon

Our first topic is to examine Typhoon in comparison to its static version and assess performance loss due to slicing and remapping. Table VII shows this result using 4-16 KB slices and breaks down the speed for individual levels  $\mathcal{L}_0, \dots, \mathcal{L}_5$ . For 32-bit keys in the first four columns, the single-threaded Typhoon shows a 1-5% advantage over the static version due to slice reuse and ties S-Typhoon at  $\mathcal{L}_3$ . Even 4-KB slices allow Typhoon to finish the sort on levels 0-4 faster than S-Typhoon (i.e., 259 vs 256M/sec). Adding remapping at  $\mathcal{L}_5$ , which is 8-13× faster than the splitter, yields a final Typhoon speed between 252 and 261M/sec, depending on slice size. This is quite competitive against S-Typhoon; in fact, slices 8 KB or larger lead to no loss of performance.

For 64-bit items (i.e., 32-bit keys with 32-bit values), the number of records per slice is reduced by half, which means that all slice-related activities occur twice as frequently on a per-key basis. Thus, it is not surprising that in these cases Typhoon needs double the slice size to achieve the same relative performance (e.g., 16 KB to match S-Typhoon on 64-bit items vs 8 KB on 32-bit, shown in bold in Table VII). For multi-threaded cases on the right side of the table, Typhoon again generally runs faster than S-Typhoon when splitting on  $\mathcal{L}_0 - \mathcal{L}_2$ ; however, it now loses up to 18% on the histogram pass (i.e., level  $\mathcal{L}_3$ ). This is because the optimized version Hv4 does not work with  $T \geq 2$  threads and the sort has to use the slower Hv3. Adding the cost of  $\mathcal{L}_5$ , where the OS struggles to maintain linear scaling of remapping speed, results in 1-5% loss on the full sort. Nevertheless, Typhoon-16KB hits 1874/2088 = 89% of `memcpy` bandwidth using 32-bit items and 976/1044 = 93% using 64-bit.

TABLE IX  
SINGLE-CORE SPEED (M/SEC) ON 32-BIT KEYS

Sort	$\mathcal{D}_1$	$\mathcal{D}_2$	$\mathcal{D}_3$	$\mathcal{D}_4$	$\mathcal{D}_5$	$\mathcal{G}$
Gorset [14]	37	52	51	38	42	44
Polychroniou [25]	34	34	32	32	30	—
Ska [29]	40	96	81	51	84	84
Regions [22]	77	58	93	79	96	85
Voracious [24]	79	80	86	81	84	86
Vortex [15]	150	122	128	135	147	127
IPS <sup>2</sup> Ra [4]	46	107	121	58	101	127
Dovetail [10]	103	99	94	103	102	99
Reinald [26]	96	100	103	101	100	111
Fast-Radix [31]	69	68	71	70	70	72
DFR [30]	76	69	97	67	79	129
pdqsort [23]	34	55	80	34	53	56
Blacher-256 [6]	133	109	117	133	133	131
IPS <sup>4</sup> o [4]	36	50	51	36	50	55
Highway-512 [13]	115	128	176	115	115	140
Intel-512 [17]	149	158	48	154	153	78
Origami-512 [3]	131	131	131	131	131	131
Typhoon-16KB	257	259	261	260	259	261
	1.7×	1.6×	1.5×	1.7×	1.7×	1.9×

Table VIII examines strong scaling of 32-bit speed as the number of threads increases, including a multiplicative factor improvement compared to the single-threaded version. We fix the input size at 1 GB and set the affinity mask to one thread per core, which yields the best result. From the table, observe that splitter speed (levels  $\mathcal{L}_0 - \mathcal{L}_2, \mathcal{L}_4$ ) scales almost perfectly until 6 cores, but then slows down to  $6.7\times$  at 7 cores and  $7.3\times$  at 8 cores as it starts approaching RAM bandwidth. At the peak, the splitter reaches 8.8B keys/sec, or 35.3 GB/s. On the other hand, the histogram at  $\mathcal{L}_3$  shows a noticeably worse scaling behavior, which arises from the fact that the single-threaded version Hv4 has an 11% advantage over the multi-threaded Hv3. It is also interesting to observe that the OS fails to linearly scale its remapping speed on  $\mathcal{L}_5$ , finishing with a  $5.2\times$  speed-up on 8 cores. Considering all these factors, Typhoon’s final speed in the bottom row is quite reasonable.

### B. Baseline Sorts

For the next group of tests, we use five synthetic 8-GB datasets: uniformly random integers ( $\mathcal{D}_1$ ); a sorted sequence of uniform numbers, where every 7-th key is set to `UINT_MAX` ( $\mathcal{D}_2$ ); uniformly random keys, each repeated  $U$  times, where  $U$  is drawn from the Zipf distribution with  $\alpha = 1, \beta = 7$ , then shuffled randomly ( $\mathcal{D}_3$ ); integer keys drawn from a normal distribution with mean `UINT32_MAX/2` and standard deviation equal to 1/3 of the mean ( $\mathcal{D}_4$ ); and uniformly random floats between 0 and `FLT_MAX` ( $\mathcal{D}_5$ ). We also use one real-world dataset, which is an inter-domain out-graph  $\mathcal{G}$  from the IRLbot web crawl [18], consisting of 89M nodes and 1.8B edges. We leverage  $\mathcal{G}$  for two standard applications – *computation of in-degree*, which entails sorting 7.2 GB of out-neighbor adjacency lists, and *graph inversion*, which requires either a stable key-value sort on 14.4 GB of (dest, src) pairs, where each node ID is 32-bit, or an unstable 64-bit key sort.

Table IX shows the speed on 32-bit keys, where we partition prior work into four groups (top to bottom) – MSD (most

TABLE X  
ALL-CORE SPEED (M/SEC) ON 32-BIT KEYS

Sort	$\mathcal{D}_1$	$\mathcal{D}_2$	$\mathcal{D}_3$	$\mathcal{D}_4$	$\mathcal{D}_5$	$\mathcal{G}$
Regions [22]	689	667	689	700	675	761
Voracious [24]	581	906	566	597	587	688
IPS <sup>2</sup> Ra [4]	526	967	1049	650	777	816
Dovetail [10]	312	350	257	339	326	267
IPS <sup>4</sup> o [4]	327	432	450	327	417	458
Origami-512 [3]	919	927	930	939	946	931
Typhoon-16KB	1,879	1,879	1,920	1,891	1,915	1,912
	2.0×	1.9×	1.8×	2.0×	2.0×	2.1×

significant digit) radix sort, LSD (least significant digit) radix sort, quick/sample sort, and merge sort – each in chronological order of publication. When a method relies on SIMD, we specify after its name the vector width (i.e., 128, 256, or 512) used in the benchmark. We highlight the fastest prior approach in each column with a gray background, run Typhoon with 16-KB slices (i.e., 12 MB of aux memory per thread), and show the speed-up factor against the best alternative in the bottom row. Dashes indicate inability to finish the sort (e.g., crashing, failing sortedness checks, unsupported input size).

There are three types of methods that stand out – the Vortex MSD [15], an SIMD quick sort from Blacher [6], with more general implementations at Google [13] and Intel [17], and the Origami SIMD merge sort [3]. All three perform quite well, delivering over 130M/sec on at least one dataset, but there is no clear winner between them. Origami posts remarkably stable speed in all columns, but it is neither the fastest nor in-place. Vortex wins on uniform data, but drops 19% between  $\mathcal{D}_1$  and  $\mathcal{D}_2$ . Blacher/Highway have similar levels of fluctuation, while Intel takes a 68% dive on  $\mathcal{D}_3$  and 48% on  $\mathcal{G}$ . In contrast, Typhoon wins in all six columns, runs in-place, and posts a 50–90% improvement over the best prior methods. Furthermore, its speed on non-uniform data is never slower than on  $\mathcal{D}_1$ , while the deviation between the max and the min is only 1.9%. This was expected from its robustness against key non-uniformity.

Among the 17 prior methods in Table IX, only six support multi-threading. Their all-core speed is displayed in Table X. Across related work, Origami wins in the uniform case by a large margin and delivers the best result in three additional columns; however, it needs double the RAM of Typhoon and uses power-hungry AVX-512 intrinsics to achieve this level of performance. In contrast, Typhoon in Table X operates mostly using scalar instructions and still almost doubles the speed of the best prior work in all six cases.

For the next experiment, we use 64-bit items composed of 32-bit key-value pairs. Some of the prior work does not have a separate provision for this case, requiring that such items be treated as monolithic 64-bit keys. Reasons include faster performance (e.g., in comparison-based SIMD methods) and ability to use unstable sorts to achieve common database tasks that would otherwise need stability (e.g., graph inversion). In this comparison, we omit Blacher [6] since it only works with 32-bit keys and add another MSD method Raduls2 [19], which

TABLE XI  
SINGLE-CORE SPEED (M/SEC) ON 64-BIT KEY-VALUE PAIRS

Sort	$\mathcal{D}_1$	$\mathcal{D}_2$	$\mathcal{D}_3$	$\mathcal{D}_4$	$\mathcal{D}_5$	$\mathcal{G}$
Gorset [14]	21	46	21	24	21	20
Polychroniou [25]	27	20	11	25	14	—
Ska [29]	36	83	68	38	67	32
Raduls2 [19]	82	65	56	92	58	53
Regions [22]	49	45	70	56	72	29
Voracious [24]	57	54	59	53	53	56
Vortex [15]	120	102	68	117	63	57
IPS <sup>2</sup> Ra [4]	45	96	104	45	88	38
Dovetail [10]	67	67	67	67	68	62
Reinald [26]	39	39	39	38	40	37
Fast-Radix [31]	40	40	43	40	43	38
DFR [30]	49	47	—	48	33	—
pdqsort [23]	31	48	32	31	31	30
IPS <sup>4</sup> o [4]	31	38	41	30	39	27
Highway-512 [13]	57	57	58	57	57	54
Intel-512 [17]	76	73	75	76	76	69
Origami-512 [3]	55	55	55	55	55	53
Typhoon-16KB	184	188	193	186	202	192
	1.5×	1.8×	1.9×	1.6×	2.3×	2.8×

was absent previously as it requires key length to be a multiple of 8 bytes.

Table XI shows the single-threaded outcome. First notice that, compared to 32-bit cases in Table IX, the three AVX-512 methods take a huge performance hit, sinking from 115-149M/sec to 55-76M/sec. Second, even though Vortex delivers excellent results for the uniform case  $\mathcal{D}_1$  (i.e., 120M/sec), it degrades to  $\sim$ 65M/sec on  $\mathcal{D}_3$  and  $\mathcal{D}_5$ . Graph inversion on  $\mathcal{G}$  also gets derailed, achieving only 57M/sec, i.e., a 2.1× reduction compared to the uniform case. This highlights the fact that real-world datasets are often skewed in a way that can heavily destabilize performance of MSD methods. In another similar case, Raduls2 suffers a 1.5× speed drop on  $\mathcal{G}$ . In contrast, Typhoon in Table XI shows consistent performance across the columns, delivering a speed-up that ranges from 1.5× on  $\mathcal{D}_1$  to 2.8× on  $\mathcal{G}$ .

Scalability to multiple threads is shown in Table XII. Among prior work, Raduls2 wins in three columns, while IPS<sup>2</sup>Ra owns the top spot for the remaining cases. As encountered before, decision between previous methods is difficult, especially considering that IPS<sup>2</sup>Ra is in-place, while Raduls2 is not. However, with the introduction of Typhoon, this choice becomes simpler – our method reliably, and by a wide margin, yields the best speed. Its current performance is stifled by insufficient RAM bandwidth, where linear scaling of the numbers in Table XI suggest a peak rate of  $184 \times 8 = 1472$ M/sec, i.e., 1.5× more than shown in Table XII.

### C. Other Platforms and In-Place Experiments

We next examine performance across seven additional CPU architectures whose characteristics are shown in Table XIII. The first three entries are server CPUs that use quad-channel memory, while the other four are dual-channel desktop chips. Sandy/Ivy Bridge implement SIMD instruction sets up to AVX, Broadwell/Coffee/Alder Lake support up to AVX2, and Zen4/Zen5 allow AVX-512. For each of the configurations,

TABLE XII  
ALL-CORE SPEED (M/SEC) ON 64-BIT KEY-VALUE PAIRS

Sort	$\mathcal{D}_1$	$\mathcal{D}_2$	$\mathcal{D}_3$	$\mathcal{D}_4$	$\mathcal{D}_5$	$\mathcal{G}$
Raduls2 [19]	656	478	394	737	433	491
Regions [22]	365	382	359	351	296	291
Voracious [24]	402	510	397	405	340	298
IPS <sup>2</sup> Ra [4]	409	737	623	418	466	318
Dovetail [10]	198	206	177	197	197	130
IPS <sup>4</sup> o [4]	286	366	351	291	341	286
Origami-512 [3]	380	389	395	394	395	392
Typhoon-16KB	986	989	998	986	1,001	997
	1.5×	1.3×	1.6×	1.3×	2.1×	2.0×

TABLE XIII  
MACHINE SPECIFICATIONS FOR IN-PLACE TESTS

Model	Year	Family	RAM	GB
Intel Xeon E5-2690	2012	Sandy Bridge (SB)	DDR3-1333	256
Intel Xeon E5-2680v2	2013	Ivy Bridge (IB)	DDR3-1866	192
Intel Xeon E5-2680v4	2016	Broadwell (BW)	DDR4-2400	128
Intel i7-8700K	2017	Coffee Lake (CL)	DDR4-3200	64
Intel i5-12600K	2021	Alder Lake (AL)	DDR5-6400	32
AMD 7950X	2022	Raphael (Zen4)	DDR5-6400	32
AMD 9600X	2024	Granite Ridge (Zen5)	DDR5-6400	32

we leave 5 GB for background processes and fill the rest with uniform keys. This leads to sort sizes 27-251 GB and requires methods that can operate in-place.

Only 10 prior implementations satisfy this criterion. Their performance on 32-bit keys is shown in Table XIV, where `std::sort` is added for reference. Note that we remove Blacher [6] since its usage of AVX2 gather instructions restricts array indexes to 4 bytes, which limits the sort to 16 GB. In the first two columns of Table XIV, Highway [13] is forced to use SSE, resulting in worse speed (i.e., 21M/sec) than a basic implementation [14] of the American Flag Sort [21] in the first row. This is in contrast to earlier Skylake-X results (Table IX), where [13] was 3.1× faster. The Intel version [17] of the same algorithm does not support SSE, while both IPSxx methods crash on inputs above  $\sim$ 128 GB. In the end, Vortex [15] squeaks out a win on Sandy Bridge and Regions [22] on Ivy Bridge, but Typhoon in the last row manages to more than double their performance.

Once AVX2 kicks in on Broadwell, SIMD methods become more competitive in the third column, with Intel climbing to the top. Vortex mounts a comeback on Coffee/Alder Lake in the next two columns, but is still 1.6-1.8× slower than Typhoon. With a jolt from AVX-512 in the last two columns, Intel almost catches up to Vortex; however, both methods are still roughly half of Typhoon’s 388M/sec on Zen4 and 491M/sec on Zen5. Similar observations hold for 64-bit key-value pairs in Table XV. Outside of Vortex, which uses up to 1 GB of extra memory and exhibits high volatility on non-uniform keys, the other prior methods are at least 2.6× slower on Coffee Lake, 3.3× on Zen4, and 3× on Zen5.

Overall, results show that across a range of desktop/server generations, Intel/AMD CPU offerings, and SSE/AVX2/AVX-512 instruction sets, Typhoon delivers the best performance, consistently taking the top spot in every comparison and

TABLE XIV  
SINGLE-CORE IN-PLACE SPEED (M/SEC) ON 32-BIT KEYS

Sort	SB	IB	BW	CL	AL	Zen4	Zen5
Gorset [14]	25	26	24	48	46	71	73
Polychroniou [25]	20	21	23	35	44	49	53
Ska [29]	40	43	41	84	99	116	120
Regions [22]	53	57	52	89	124	121	142
Vortex [15]	54	56	53	162	178	203	265
IPS <sup>2</sup> Ra [4]	—	—	47	90	109	110	121
pdqsort [23]	23	24	24	33	40	45	47
IPS <sup>4</sup> o [4]	—	—	22	33	34	46	50
Highway [13]	21	22	42	77	106	149	185
Intel [17]	—	—	63	118	167	177	240
std::sort	7	7	7	9	10	13	13
Typhoon-16KB	120 2.2×	129 2.3×	129 2.0×	265 1.6×	328 1.8×	388 1.9×	491 1.8×

TABLE XV  
SINGLE-CORE IN-PLACE SPEED (M/SEC) ON 64-BIT KEY-VALUE PAIRS

Sort	SB	IB	BW	CL	AL	Zen4	Zen5
Gorset [14]	16	16	15	32	28	47	48
Polychroniou [25]	14	15	15	27	30	33	34
Ska [29]	31	32	33	66	71	80	79
Regions [22]	36	39	47	74	84	96	99
Vortex [15]	29	31	31	133	175	187	239
IPS <sup>2</sup> Ra [4]	—	—	36	77	80	85	88
pdqsort [23]	17	19	20	29	41	47	47
IPS <sup>4</sup> o [4]	—	—	19	31	34	44	47
Highway [13]	10	11	18	35	46	92	123
Intel [17]	—	—	22	43	60	89	135
std::sort	7	7	7	9	10	13	13
Typhoon-16KB	76 2.1×	79 2.0×	83 1.8×	197 1.5×	233 1.3×	321 1.7×	404 1.7×

finishing 32-bit sorts 38× faster than std::sort on AMD Zen5. Furthermore, it is highly skew-resilient, as well as the only method in this comparison that is both stable and in-place.

## V. ACKNOWLEDGMENTS

The authors are grateful to Benjamin Ramon and Carson Hanel for initial prototyping and insightful discussion.

## VI. CONCLUSION

We developed a novel LSD sort called Typhoon and demonstrated that it worked remarkably fast across a variety of CPU architectures, memory configurations, single/multi-core scenarios, input skew, and array sizes. Not only that, but Typhoon is also stable, in-place, and distribution-insensitive. Future work involves handling longer keys and testing Typhoon in big-data frameworks/databases.

## REFERENCES

- [1] Apache Hadoop. [Online]. Available: <http://hadoop.apache.org/>.
- [2] Apache Spark. [Online]. Available: <https://spark.apache.org/>.
- [3] A. Arman and D. Loguinov, “Origami: A High-Performance Mergesort Framework,” in *Proc. VLDB*, Sep. 2022, pp. 259–271.
- [4] M. Axtmann, S. Witt, D. Ferizovic, and P. Sanders, “Engineering In-place (Shared-memory) Sorting Algorithms,” *ACM Transactions on Parallel Computing*, vol. 9, no. 1, pp. 1–62, Jan. 2022.
- [5] C. Balkesen, G. Alonso, J. Teubner, and M. T. Özsü, “Multi-core, Main-memory Joins: Sort vs. Hash Revisited,” *VLDB Endow.*, vol. 7, no. 1, pp. 85–96, Sep. 2013.
- [6] M. Blacher, J. Giesen, and L. Kühne, “Fast and Robust Vectorized In-Place Sorting of Primitive Types,” in *Proc. International Symposium on Experimental Algorithms (SEA)*, Jun. 2021, pp. 3:1–3:16.
- [7] M. Cho, D. Brand, R. Bordawekar, U. Finkler, V. Kulandaisamy, and R. Puri, “PARADIS: An Efficient Parallel Algorithm for In-place Radix Sort,” *VLDB Endow.*, vol. 8, no. 12, pp. 1518–1529, Aug. 2015.
- [8] Y. Cui, D. Xiao, D. B. Cline, and D. Loguinov, “Improving I/O Complexity of Triangle Enumeration,” *IEEE Transactions on Knowledge and Data Engineering*, vol. 34, no. 4, pp. 1815–1828, Apr. 2022.
- [9] J. Dean and S. Ghemawat, “MapReduce: Simplified Data Processing on Large Clusters,” in *Proc. USENIX OSDI*, Dec. 2004, pp. 137–150.
- [10] X. Dong, L. Dhulipala, Y. Gu, and Y. Sun, “Parallel Integer Sort: Theory and Practice,” in *Proc. ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming (PPoPP)*, Mar. 2024, pp. 301–315.
- [11] T. Downs, “Memory Disambiguation on Skylake,” 2021. [Online]. Available: <https://github.com/travisdowns/uarch-bench/wiki/Memory-Disambiguation-on-Skylake>.
- [12] K. Fukazawa and R. Takahashi, “Performance Evaluation of the Fourth-Generation Xeon with Different Memory Characteristics,” in *Proc. HPC Asia Workshops*, Jan 2024, pp. 55–62.
- [13] Google, “Highway Sort,” 2024. [Online]. Available: <https://github.com/google/highway/tree/master/hwy/contrib/sort>.
- [14] E. Gorset, “In-place Radix Sort,” Apr. 2011. [Online]. Available: <https://github.com/gorset/radix>.
- [15] C. Hanel, A. Arman, D. Xiao, J. Keech, and D. Loguinov, “Vortex: Extreme-Performance Memory Abstractions for Data-Intensive Streaming Applications,” in *Proc. ACM ASPLOS*, Mar. 2020, pp. 623–638.
- [16] Intel Corporation, “Intel 64 and IA-32 Architectures Optimization Reference Manual Volume 1.” [Online]. Available: <https://www.intel.com/content/www/us/en/developer/articles/technical/intel-sdm.html>.
- [17] Intel Corporation, “x86-simd-sort,” 2024. [Online]. Available: <https://github.com/intel/x86-simd-sort>.
- [18] IRL Web Crawling Datasets. [Online]. Available: <http://irl.cs.tamu.edu/projects/web/>.
- [19] M. Kokot, S. Deorowicz, and A. Debudaj-Grabysz, “Even Faster Sorting of (Not Only) Integers,” in *Proc. International Conference on Man-Machine Interactions*, Oct. 2017, pp. 481–491.
- [20] Z. Liu, A. Arman, and D. Loguinov, “Typhoon: A Slice-Scrambled In-Place LSD Sort,” Texas AM University, Tech. Rep., Nov. 2025. [Online]. Available: <http://irl.cs.tamu.edu/publications>.
- [21] P. M. McIlroy, K. Bostic, and M. D. McIlroy, “Engineering Radix Sort,” *Computing Systems*, vol. 6, no. 1, pp. 5–27, 1993.
- [22] O. Obeya, E. Kahssay, E. Fan, and J. Shun, “Theoretically-efficient and Practical Parallel In-place Radix Sorting,” in *Proc. ACM SPAA*, Jun. 2019, pp. 213–224.
- [23] O. Peters, “Pattern-defeating Quicksort (pdqsort),” 2021. [Online]. Available: <https://github.com/orlp/pdqsort>.
- [24] A. Piot, “Voracious Sort,” 2020. [Online]. Available: [https://github.com/lakwet/voracious\\_sort](https://github.com/lakwet/voracious_sort).
- [25] O. Polychroniou and K. A. Ross, “A comprehensive study of main-memory partitioning and its application to large-scale comparison- and radix-sort,” in *Proc. ACM SIGMOD*, Jun. 2014, pp. 755–766.
- [26] A. Reinald, P. Harris, R. Rohrer, and J. Dirk. Radix sort. [Online]. Available: [http://www.cubic.org/docs/download/radix\\_ar\\_2011.cpp](http://www.cubic.org/docs/download/radix_ar_2011.cpp).
- [27] N. Satish, C. Kim, J. Chhugani, A. Nguyen, V. Lee, D. Kim, and P. Dubey, “Fast sort on CPUs and GPUs: A case for bandwidth oblivious SIMD sort,” in *Proc. ACM SIGMOD*, Jun. 2010, pp. 351–362.
- [28] F. M. Schuhknecht, P. Khanchandani, and J. Dittrich, “On the Surprising Difficulty of Simple Things: The Case of Radix Partitioning,” *VLDB Endow.*, vol. 8, no. 9, pp. 934–937, May 2015.
- [29] M. Skarupke, “I Wrote a Faster Sorting Algorithm,” 2016. [Online]. Available: <https://probablydance.com/2016/12/27/i-wrote-a-faster-sorting-algorithm/>.
- [30] S. Thiel, “Implementing the Diverting Fast Radix Algorithm,” Aug. 2022. [Online]. Available: <https://arxiv.org/abs/2207.14334>.
- [31] S. Thiel, G. Butler, and L. Thiel, “Improving GraphChi for Large Graph Processing: Fast Radix Sort in Pre-Processing,” in *Proc. ACM IDEAS*, Jul. 2016, pp. 135–141.
- [32] Typhoon. [Online]. Available: <http://irl.cs.tamu.edu/projectsstreams>.
- [33] J. Wassenberg and P. Sanders, “Engineering a Multi-core Radix Sort,” in *Proc. Euro-Par*, Aug. 2011, pp. 160–169.
- [34] H. Wong, “Store-to-Load Forwarding and Memory Disambiguation in x86 Processors,” 2014. [Online]. Available: <https://blog.stuffedcow.net/2014/01/x86-memory-disambiguation/>.